


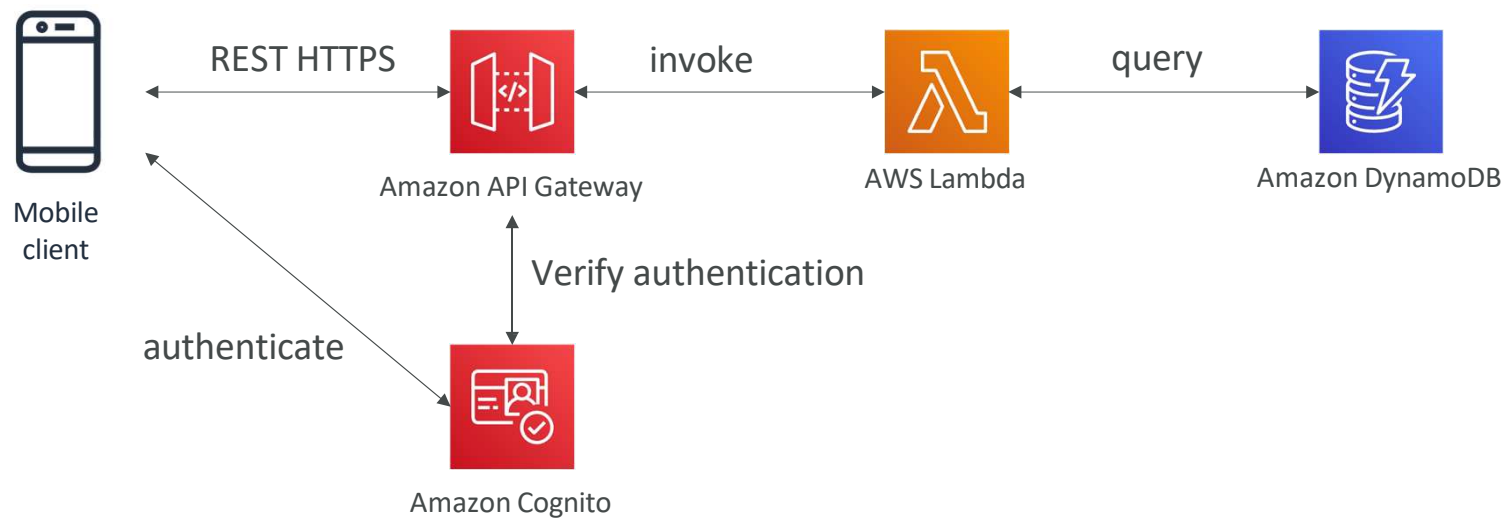
Serverless Architectures



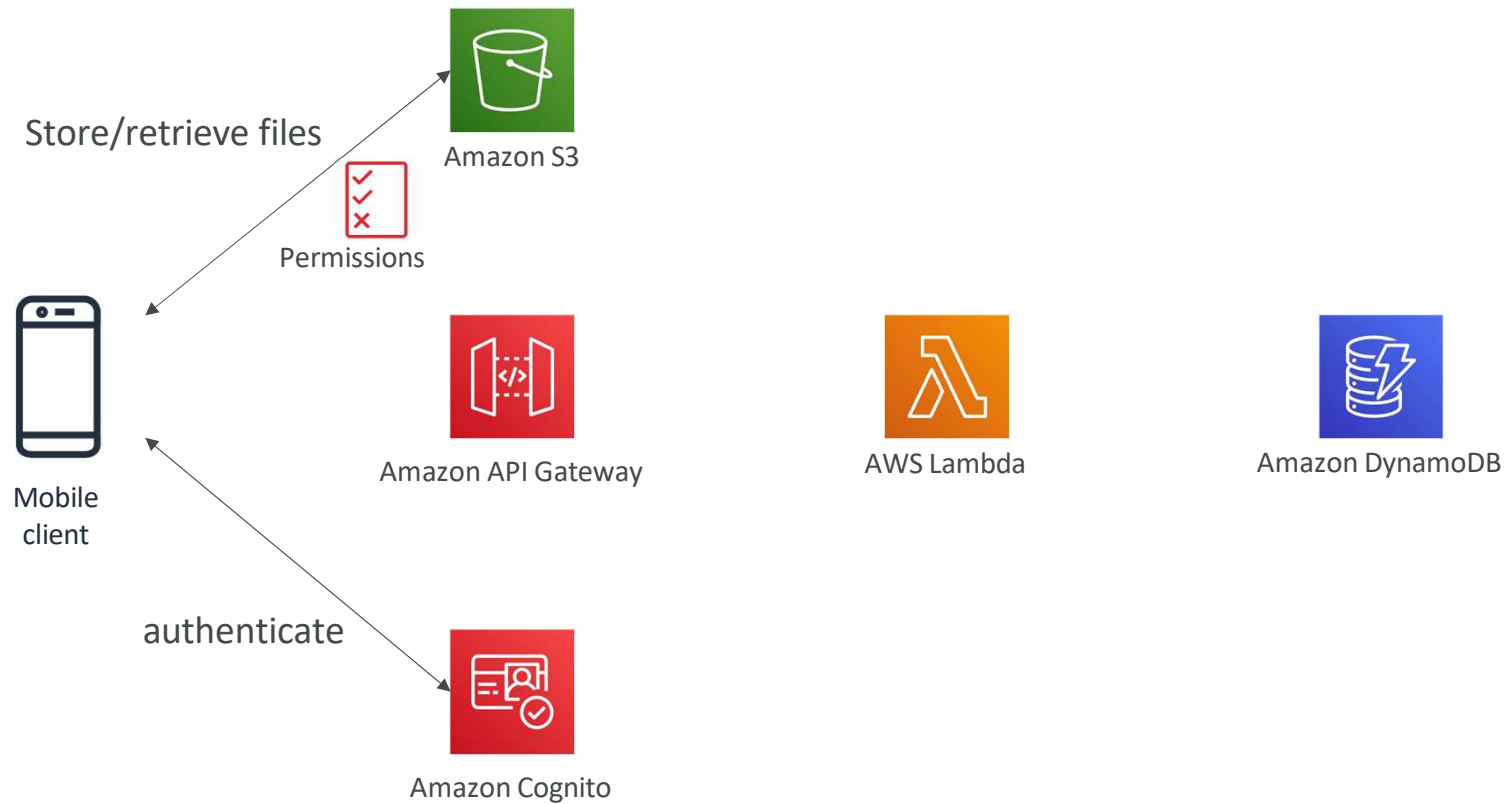
Mobile application: MyToDoList

- We want to create a mobile application with the following requirements
 - Expose as REST API with HTTPS
 - Serverless architecture
 - Users should be able to directly interact with their own folder in S3
 - Users should authenticate through a managed serverless service
 - The users can write and read to-dos, but they mostly read them
 - The database should scale, and have some high read throughput
- 

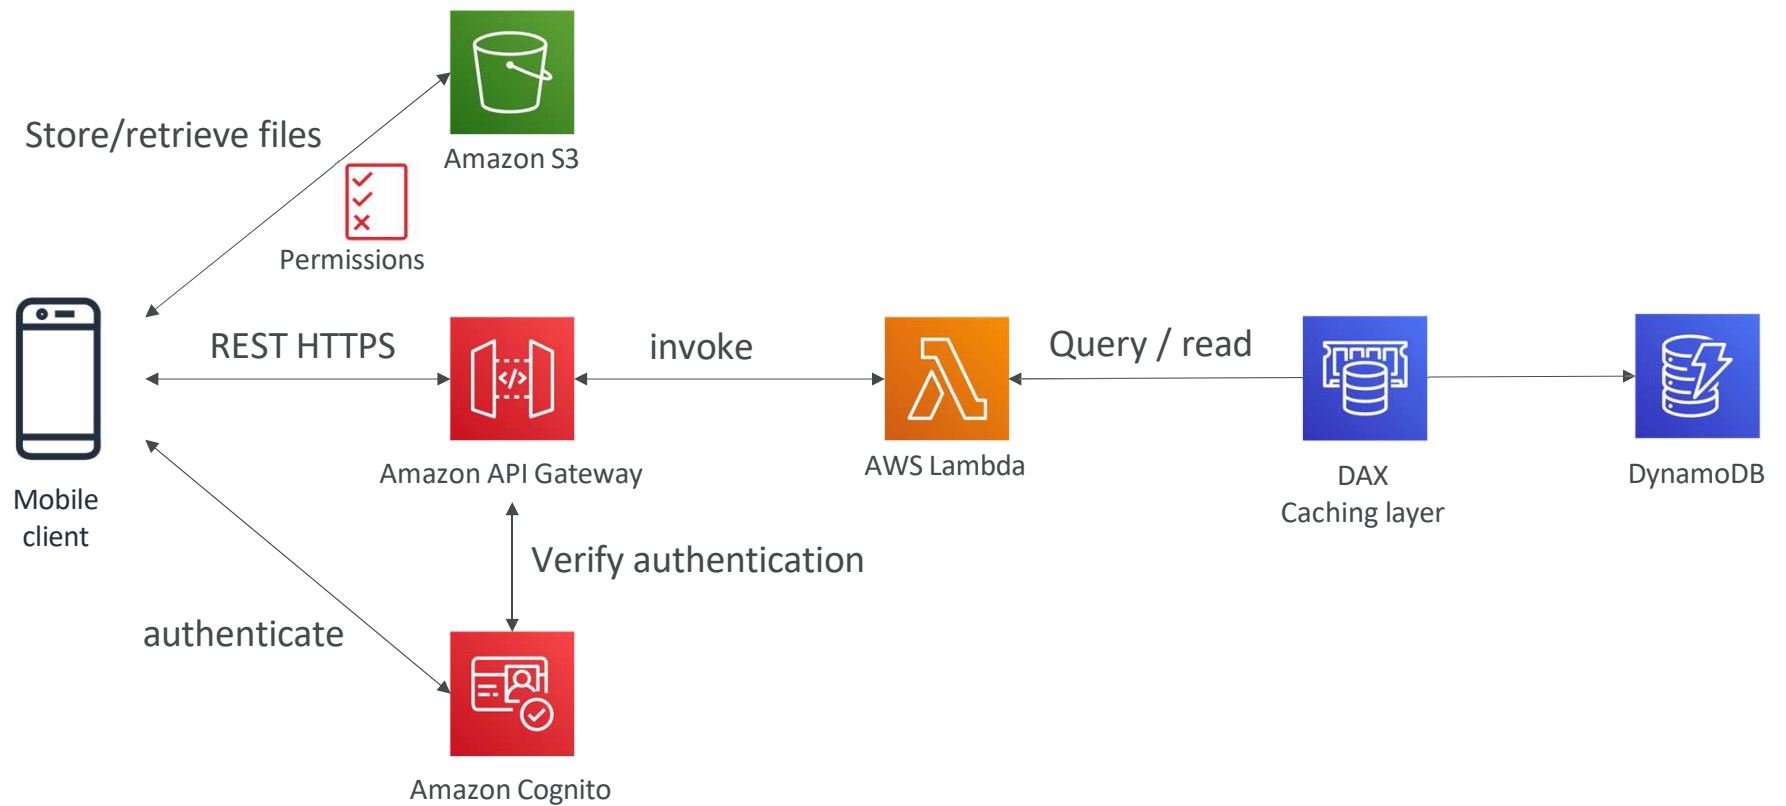
Mobile app: REST API layer



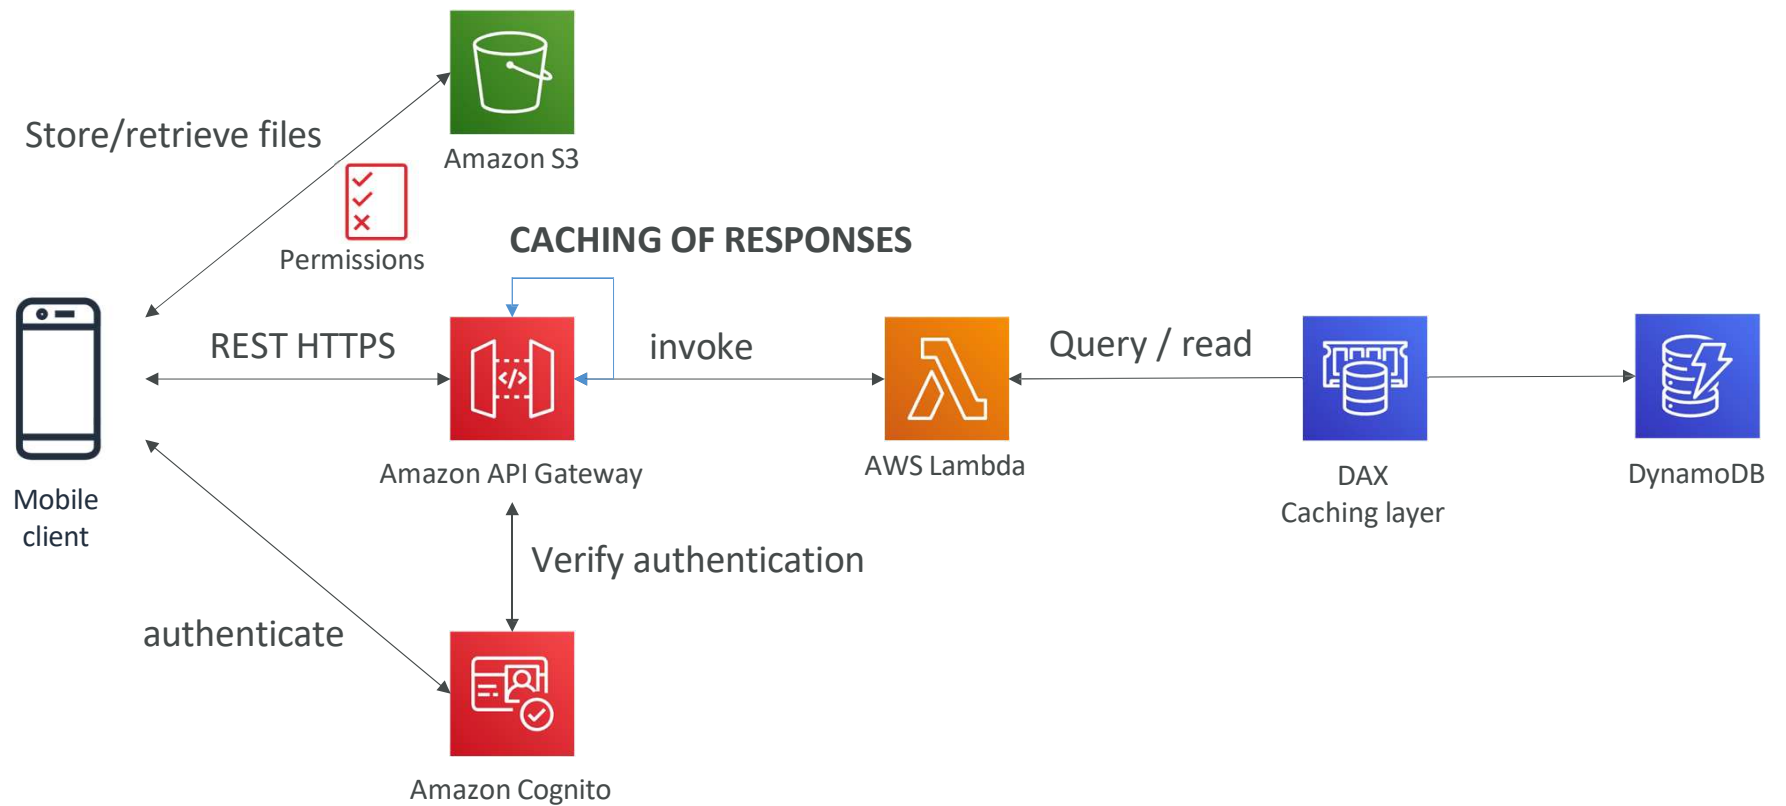
Mobile app: giving users access to S3




Mobile app: high read throughput, static data




Mobile app: caching at the API Gateway



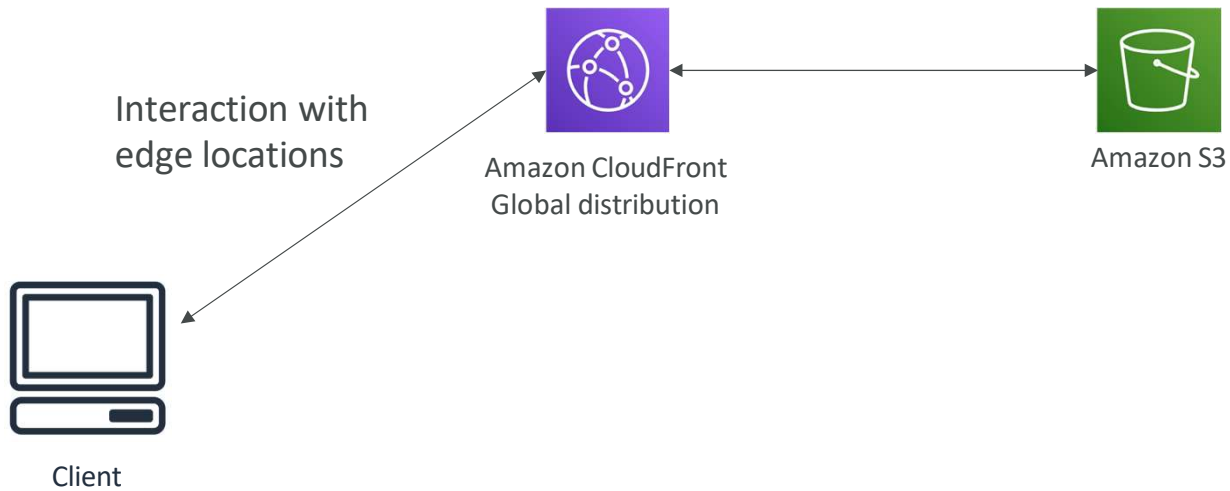
Summary

- Serverless REST API: HTTPS, API Gateway, Lambda, DynamoDB
 - Using Cognito to generate temporary credentials to access S3 bucket with restricted policy. App users can directly access AWS resources this way. Pattern can be applied to DynamoDB, Lambda...
 - Caching the reads on DynamoDB using DAX
 - Caching the REST requests at the API Gateway level
 - Security for authentication and authorization with Cognito
- 

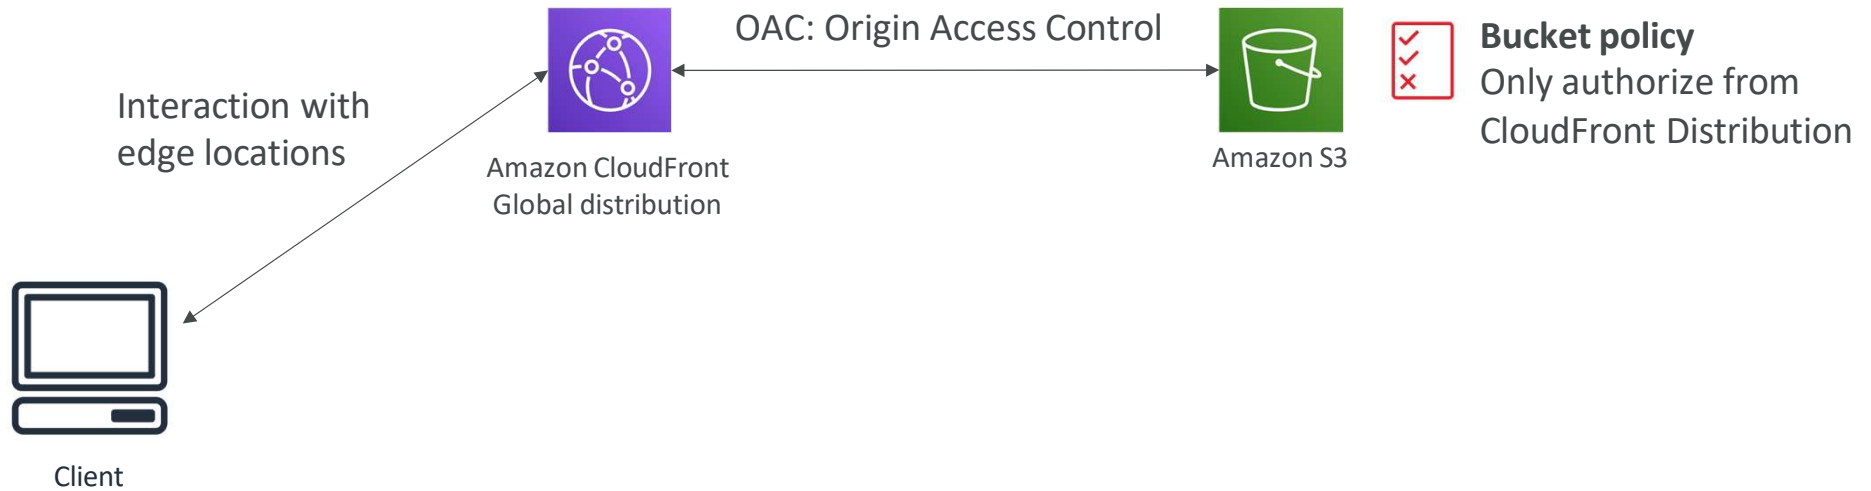
Serverless hosted website: MyBlog.com

- This website should scale globally
 - Blogs are rarely written, but often read
 - Some of the website is purely static files, the rest is a dynamic REST API
 - Caching must be implement where possible
 - Any new users that subscribes should receive a welcome email
 - Any photo uploaded to the blog should have a thumbnail generated
- 

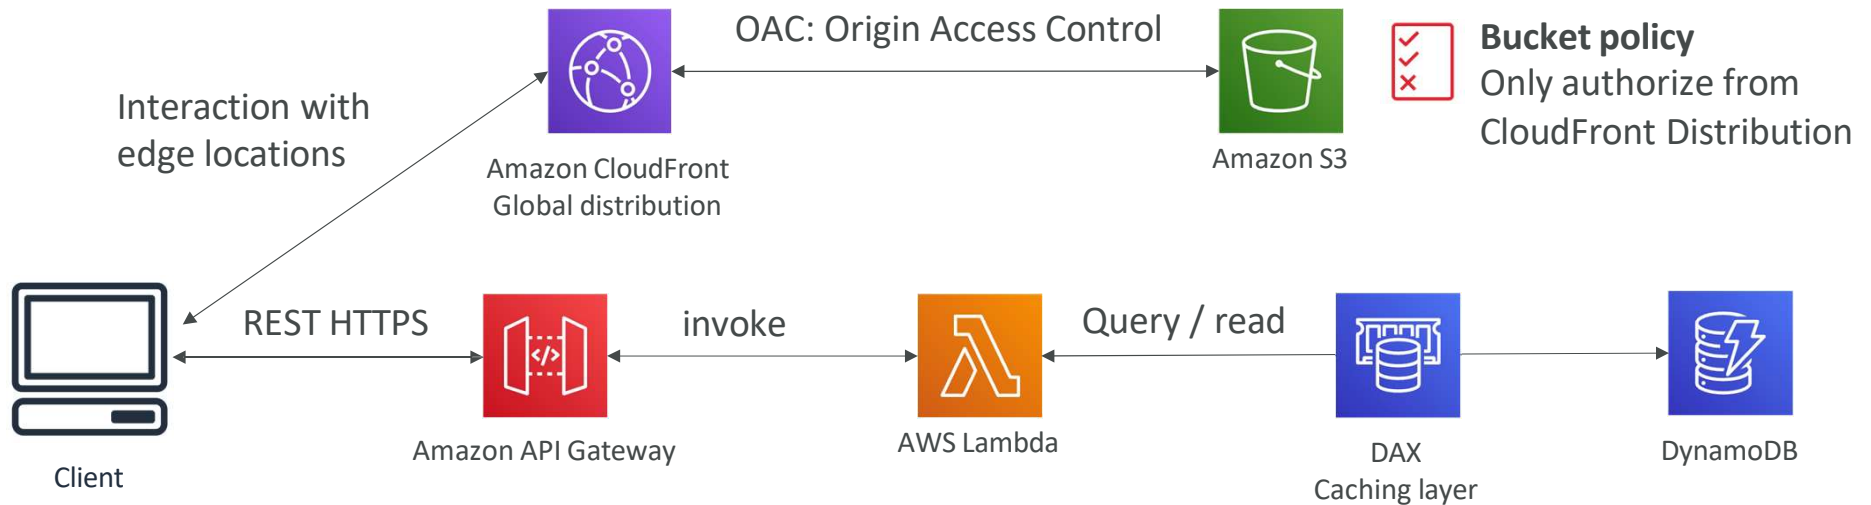
Serving static content, globally



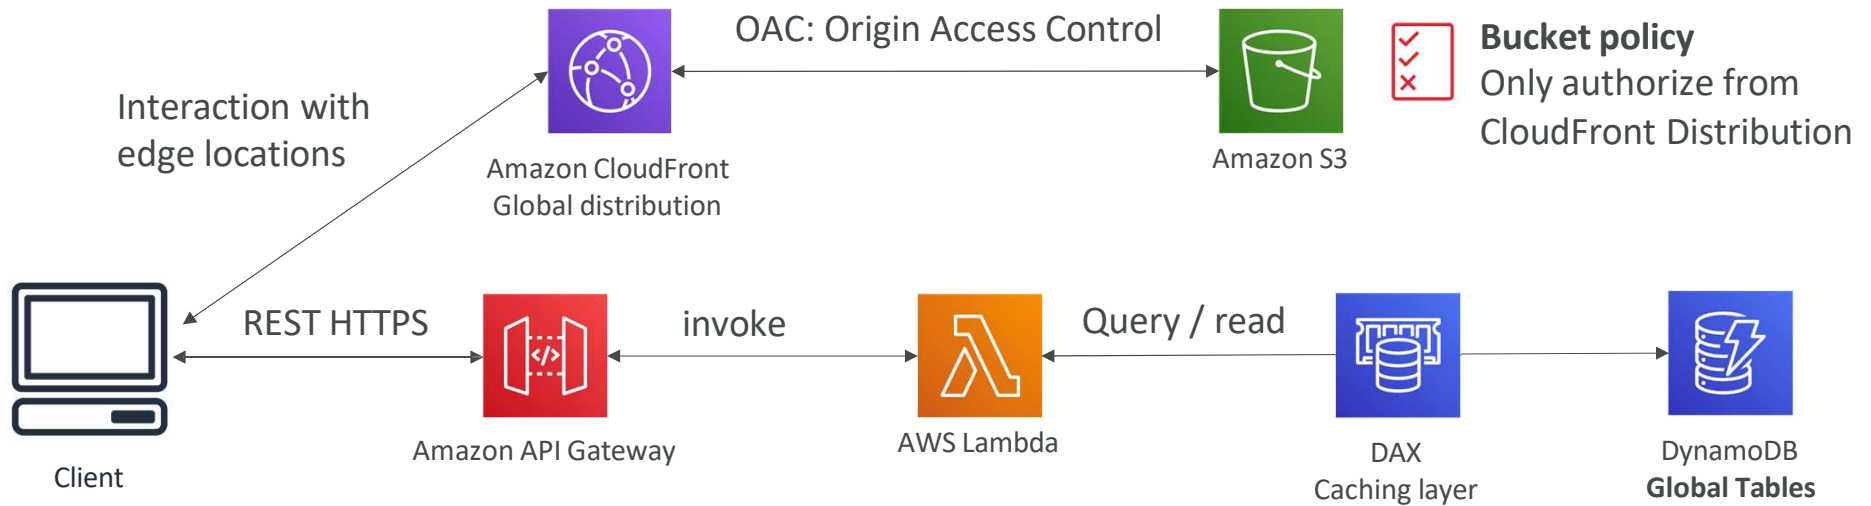
Serving static content, globally, securely



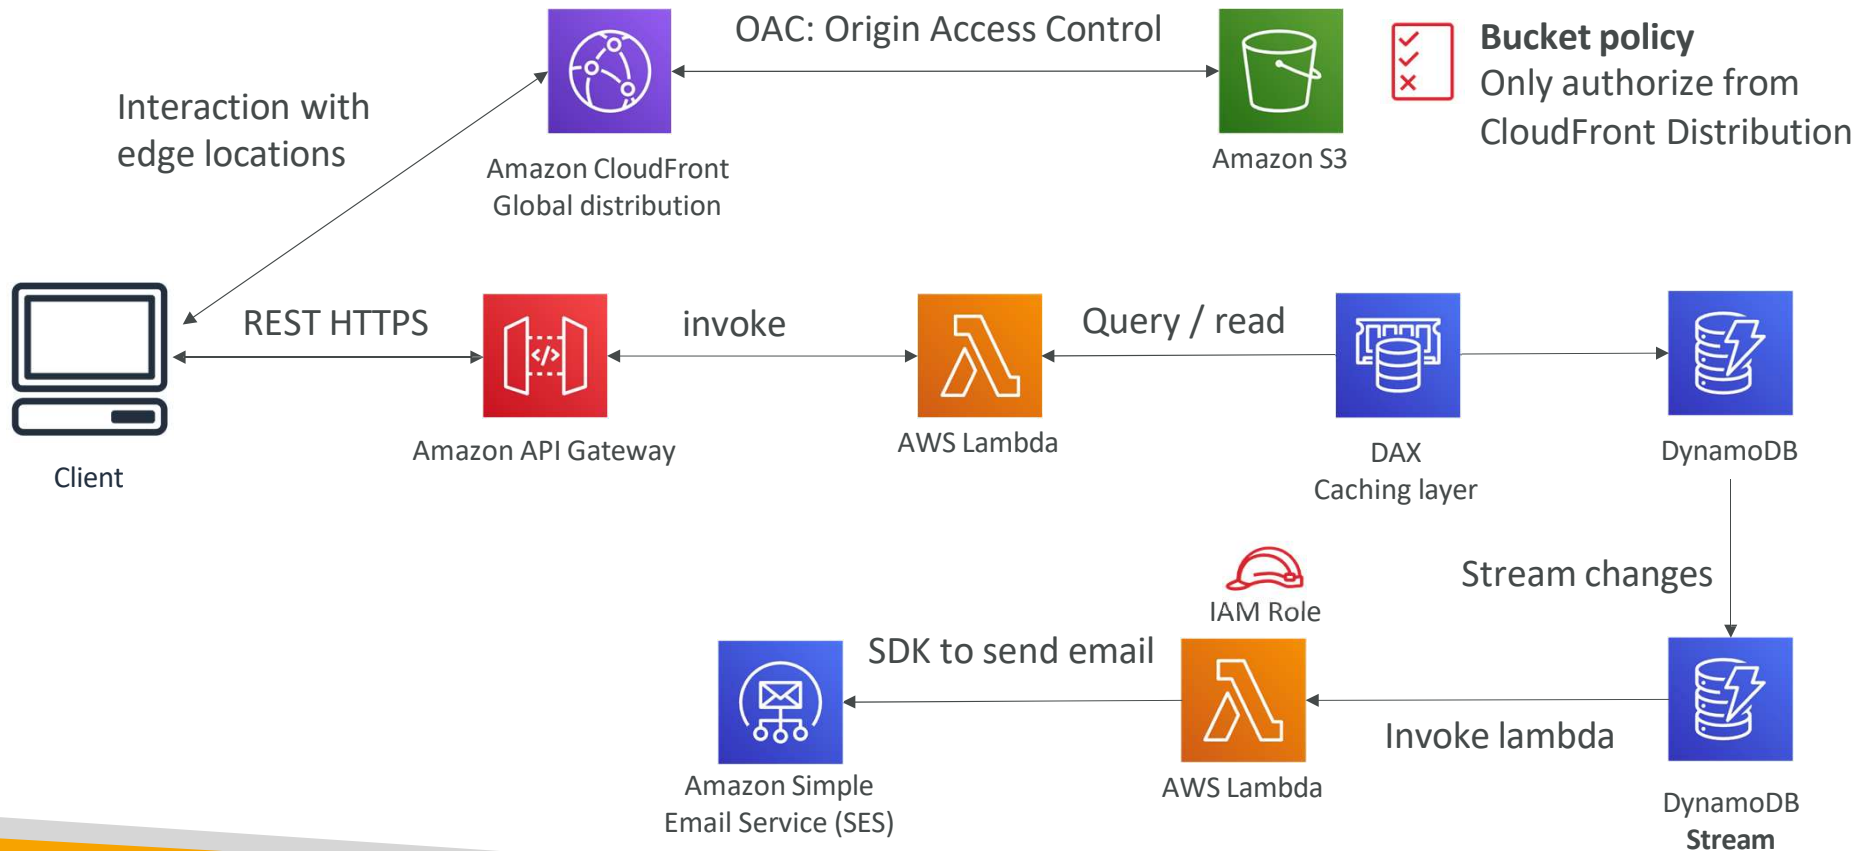
Adding a public serverless REST API



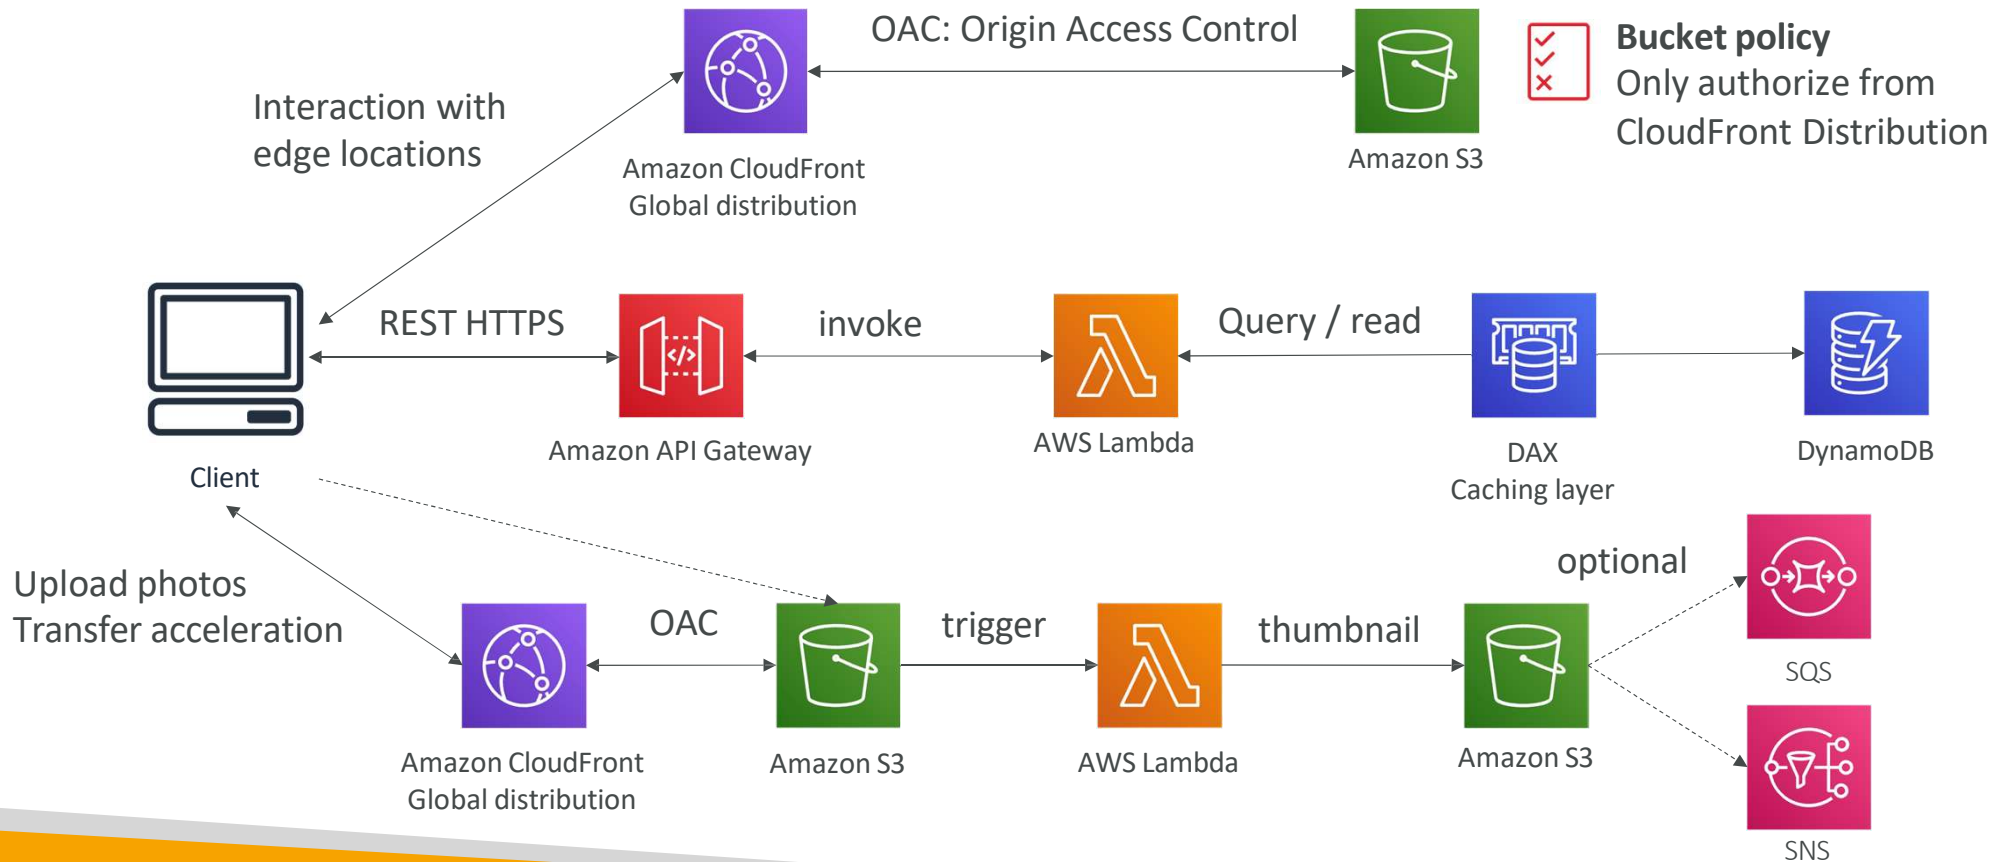
Leveraging DynamoDB Global Tables




User Welcome email flow



Thumbnail Generation flow



AWS Hosted Website Summary

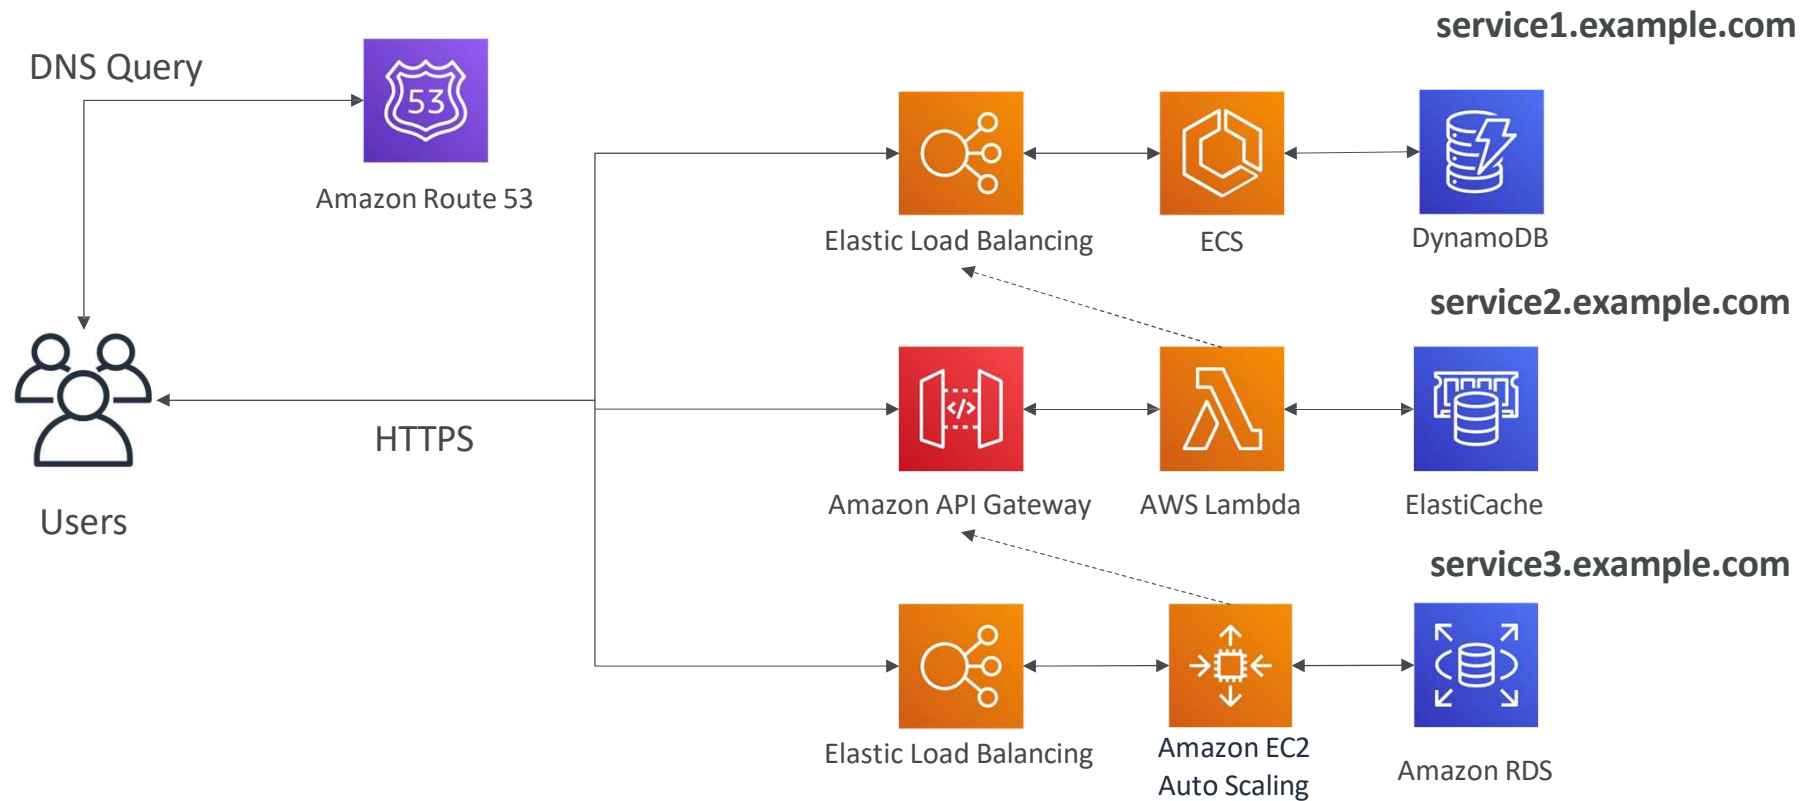
- We've seen static content being distributed using CloudFront with S3
 - The REST API was serverless, didn't need Cognito because public
 - We leveraged a Global DynamoDB table to serve the data globally
 - (we could have used Aurora Global Database)
 - We enabled DynamoDB streams to trigger a Lambda function
 - The lambda function had an IAM role which could use SES
 - SES (Simple Email Service) was used to send emails in a serverless way
 - S3 can trigger SQS / SNS / Lambda to notify of events
- 

Micro Services architecture


- We want to switch to a micro service architecture
- Many services interact with each other directly using a REST API
- Each architecture for each micro service may vary in form and shape
- We want a micro-service architecture so we can have a leaner development lifecycle for each service



Micro Services Environment



Summary on Micro Services

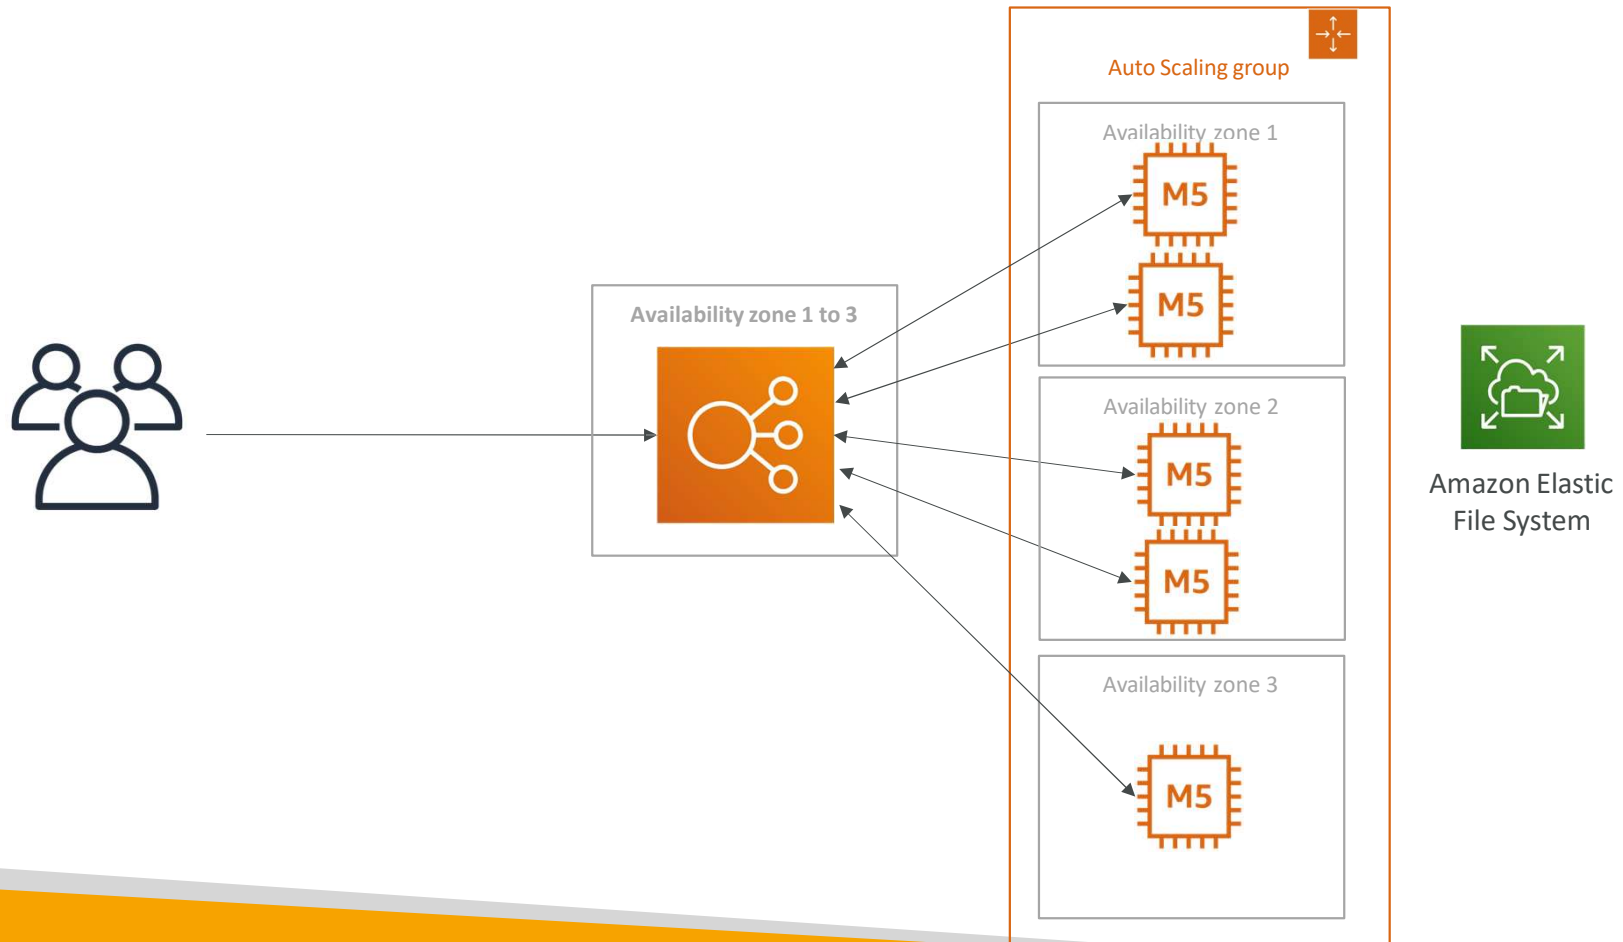
- You are free to design each micro-service the way you want
 - Synchronous patterns: API Gateway, Load Balancers
 - Asynchronous patterns: SQS, Kinesis, SNS, Lambda triggers (S3)
 - Challenges with micro-services:
 - repeated overhead for creating each new microservice,
 - issues with optimizing server density/utilization
 - complexity of running multiple versions of multiple microservices simultaneously
 - proliferation of client-side code requirements to integrate with many separate services.
 - Some of the challenges are solved by Serverless patterns:
 - API Gateway, Lambda scale automatically and you pay per usage
 - You can easily clone API, reproduce environments
 - Generated client SDK through Swagger integration for the API Gateway
- 

Software updates offloading

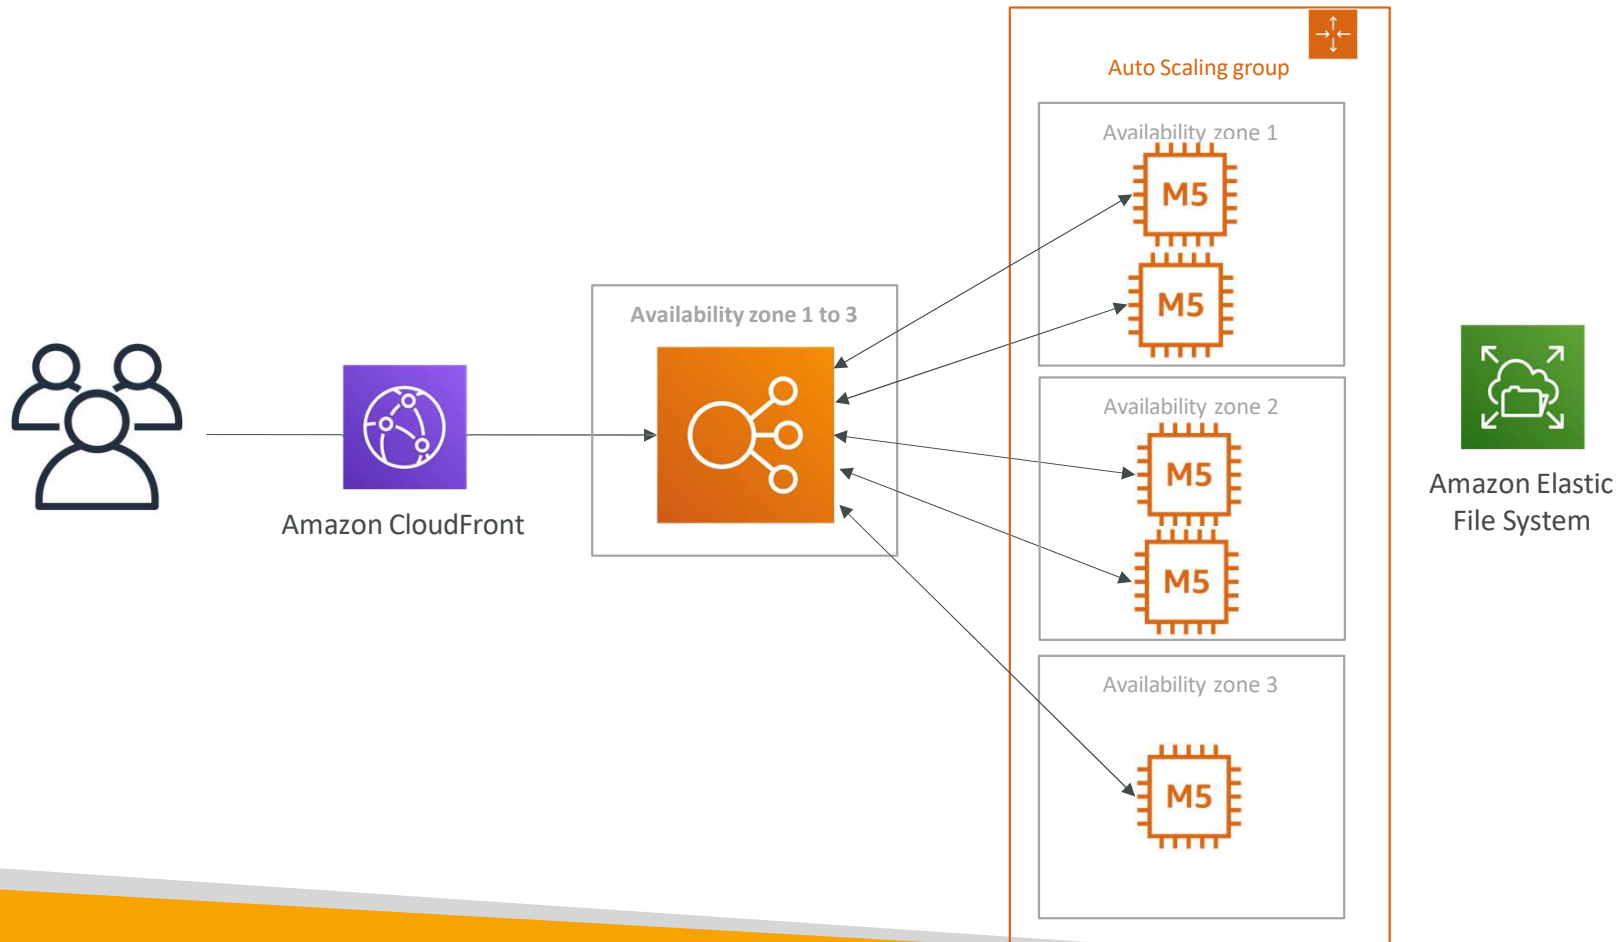
- We have an application running on EC2, that distributes software updates once in a while
- When a new software update is out, we get a lot of request and the content is distributed in mass over the network. It's very costly
- We don't want to change our application, but want to optimize our cost and CPU, how can we do it?



Our application current state



Easy way to fix things!



Why CloudFront?

- No changes to architecture
 - Will cache software update files at the edge
 - Software update files are not dynamic, they're static (never changing)
 - Our EC2 instances aren't serverless
 - But CloudFront is, and will scale for us
 - Our ASG will not scale as much, and we'll save tremendously in EC2
 - We'll also save in availability, network bandwidth cost, etc
 - Easy way to make an existing application more scalable and cheaper!
- 